

Assignment 1

ZHANG YIFEI

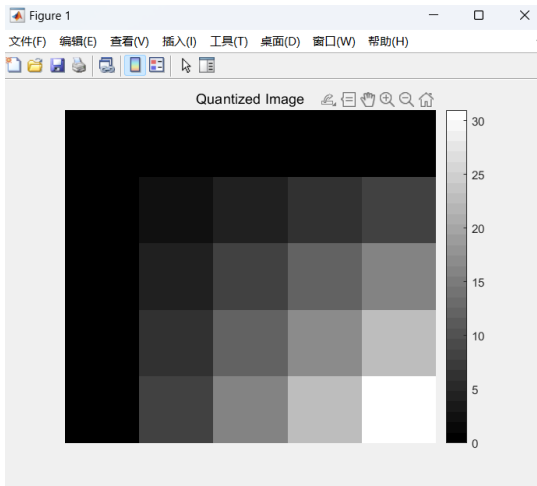
1 Image Sampling

First, create the grid and assign coordinate values to each grid.

Second, calculate the corresponding intensity values according to the formula.

Calculate the gray value based on gray level and intensity value.

```
1 % Define the resolution of the discrete image
2 num_pixels = 5;
3 % Define the number of gray levels
4 num_gray_levels = 32;
5
6 % Create a meshgrid for the x and y values
7 [x, y] = meshgrid(linspace(0, 1, num_pixels), linspace(1, 0,
8 num_pixels));
9
10 % Calculate the intensity values for each pixel using the
11 given function
12 intensity = x .* (1 - y);
13
14 % Quantize the intensity values to 32 gray levels
15 quantized_intensity = round(intensity * (num_gray_levels -
16 1));
17
18 % Create a colormap for the 32 gray levels
19 colormap(gray(num_gray_levels));
20
21 % Display the discrete image
22 imagesc(quantized_intensity);
23 axis off;
24 title('Quantized Image');
25 colorbar;
```



0	0	0	0	0
0	2	4	6	8
0	4	8	12	16
0	6	12	17	23
0	8	16	23	31

2 Histogram Equalization

$$S = T(r) = \int_0^r p_r(r) dr = 3r^2 - 2r^3$$

```

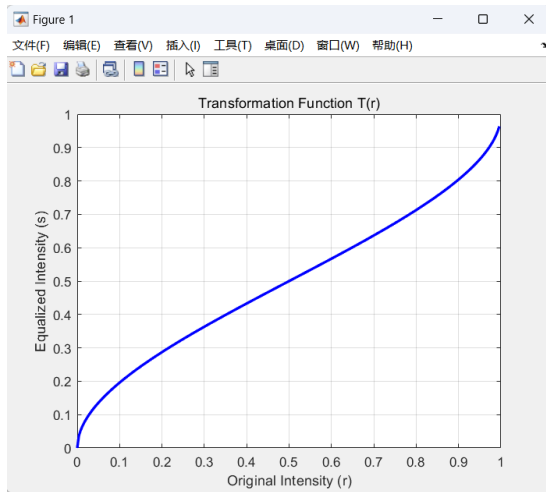
1 % Define the range
2 r = linspace(0, 1, 256);
3
4 % Define pr
5 pr = 6 * r .* (1 - r);
6
7 % Calculate F(r)
8 F_r = cumtrapz(r, pr);
9
10 % Calculate the inverse of the F(r) to get the transformation
    s = T(r)
11 F_inverse = interp1(F_r, r, linspace(0, 1, numel(r)));
12
13 % Create a figure to visualize the transformation matrix T(r)

```

```

14 figure;
15 plot(r, F_inverse, 'b', 'Linewidth', 2);
16 xlabel('Original Intensity (r)');
17 ylabel('Equalized Intensity (s)');
18 title('Transformation Function T(r)');
19 grid on;
20

```



3 Neighborhood of Pixels

```

1 warning('off','all')
2 warning
3
4 % Define your binary image (replace this with your binary
  image data)
5 binary_image = [
6     3 3 3 3 3 2 2 2 2 2 3 3 3 3 3;
7     3 3 1 2 1 0 0 0 0 0 1 2 3 3 3;
8     3 3 2 0 0 0 0 0 0 0 0 0 2 3 3;
9     3 2 0 0 1 0 0 0 0 0 1 0 0 2 3;
10    3 1 0 2 3 1 0 0 0 3 3 1 0 1 3;
11    2 0 0 3 3 2 0 0 1 3 3 2 0 0 2;
12    2 0 0 2 3 1 0 0 0 2 3 1 0 0 2;
13    2 0 0 0 0 0 0 0 0 0 0 0 0 0 2;
14    2 0 0 0 0 0 0 0 0 0 0 0 0 0 2;
15    2 0 1 2 1 1 0 0 0 1 1 2 2 0 2;
16    3 1 0 2 3 3 3 3 3 3 3 3 0 1 3;
17    3 2 0 0 2 3 3 3 3 3 2 0 0 2 3;
18    3 3 2 0 0 2 3 3 3 2 0 0 2 3 3;
19    3 3 3 2 1 0 0 0 0 0 1 2 3 3 3;

```

```

20     3 3 3 3 3 2 2 2 2 2 3 3 3 3;
21 ];
22
23 pixel = 1;
24 % Define 8-connected neighbors
25 neighbors = [-1, -1; -1, 0; -1, 1; 0, 1; 1, 1; 1, 0; 1, -1; 0,
26             -1];
27
28 % Initialize labeled image and label counter
29 global labeled_visited_coordinates
30 global labeled_img
31 global connected_components
32 labeled_visited_coordinates = zeros(size(binary_image));
33 labeled_img = zeros(size(binary_image));
34 current_label = 0;
35 connected_components = cell(0);
36 %% Threshold the image
37 for i = 1:size(binary_image, 1)
38     for j = 1:size(binary_image, 2)
39         if binary_image(i, j) <= 1
40             binary_image(i, j) = 0;
41         else
42             binary_image(i, j) = 1;
43         end
44     end
45 end
46 %%
47 % Iterate through the binary image to find and label connected
48 % components
49 for i = 1:size(binary_image, 1)
50     for j = 1:size(binary_image, 2)
51         if labeled_visited_coordinates(i, j) == 0
52             if binary_image(i, j) == pixel
53                 current_label = current_label + 1;
54                 connected_components{current_label} = [];
55                 dfs(pixel, i,
56                    j, current_label, binary_image, labeled_img,
57                    labeled_visited_coordinates, neighbors, connected_components);
58             end
59             labeled_visited_coordinates(i, j) = 1;
60         end
61     end
62 end
63 % Display the labeled image
64 imshow(label2rgb(labeled_img, 'jet', 'k'),
65         'InitialMagnification', 'fit');
66 title('Connected Components (8-connected)');
67 disp(labeled_img);

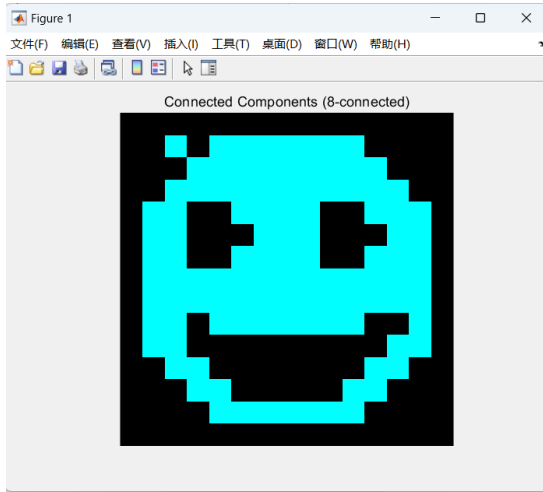
```

```

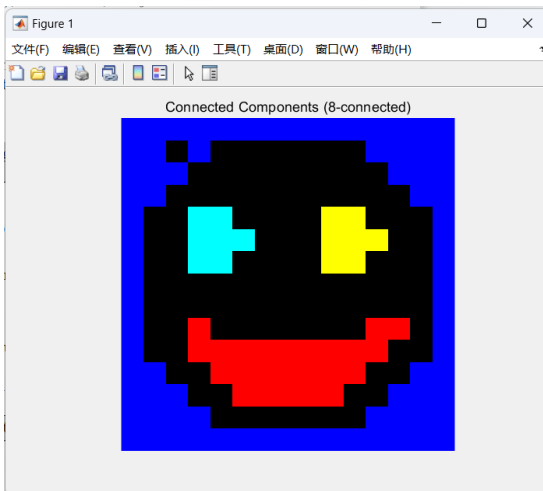
1  % Function to perform DFS for connected components labeling
2  function dfs(pixel,i,
3      j,current_label,binary_image,labeled_img,
4      labeled_visited_coordinates,neighbors,connected_components)
5      global labeled_visited_coordinates
6      global labeled_img
7      global connected_components
8
9      if labeled_visited_coordinates(i, j) == 0
10         labeled_visited_coordinates(i, j) = 1;
11         if binary_image(i,j)==pixel
12             labeled_img(i, j) = current_label;
13             connected_components{current_label} =
[connected_components{current_label}; [i, j]];
14
15             % Recursively call DFS on neighboring pixels
16             for k = 1:size(neighbors, 1)
17                 ni = i + neighbors(k, 1);
18                 nj = j + neighbors(k, 2);
19                 if ni < 1 || ni > size(binary_image, 1) || nj
< 1 || nj > size(binary_image, 2)
20                     continue;
21                 else
22                     dfs(pixel,ni,nj,current_label,binary_image,labeled_img,
23                         labeled_visited_coordinates,neighbors,connected_components);
24                 end
25             end
26         end
27     end
28 end

```

Elements that has intensity 0 or 1 (non-black part):



8-connected components for $g = 1$ (non-black part):



4 Segmentation Part of OCR

```

1 function S = im2segment(img)
2 img = uint8(img);
3 % figure;
4 % imshow(img);
5
6 % Specify the standard deviation of the Gaussian filter (to
  control the degree of blurring)
7 sigma = 0.5;
8
9 % Perform Gaussian filtering

```

```

10  img = imgaussfilt(img, sigma);
11
12  % img = imbilatfilt(img);
13  % img = medfilt2(img, [3, 3])
14
15  threshold = 0.158 % Set the threshold of image binarize
16  binary_image = imbinarize(img, threshold); % binarize
17  % imshow(binary_image);
18
19  %% 8-connected components
20  labeledImage = bwlabel(binary_image, 8);
21  minPixels = 1 ; % set min pixel num
22
23
24  %% Extracting information about connected components
25  stats = regionprops(labeledImage, 'BoundingBox',
26  'PixelIdxList');
27
28  % Initialize an array of cells for storing segmented images
29  numStats = numel(stats);
30  S = cell(1, numStats);
31
32  % Create an array of flags to keep track of merged cells
33  merged = false(1, numStats);
34  small = false(1, numStats);
35
36  %% Storing coordinate indexes of different labels in cells
37  for i = 1:numStats
38      % Get the coordinate index of the current connected
39      component
40      pixelIdxList = stats(i).PixelIdxList;
41
42      % Create a segmented image of the same size as the
43      original image
44      segmented_image = zeros(size(labeledImage));
45      segmented_image(pixelIdxList) = 1;
46      numPixels = sum(segmented_image(:) == 1);
47      if numPixels < minPixels
48          small(i)=true;
49      end
50      % Storing Segmented Images into Cells
51      S{i} = segmented_image;
52  end
53
54  %%
55  Dis_threshold = 20; % distance threshold
56
57  % Calculate the center coordinates of each cell and combine
58  them
59  for i = 1:numStats

```

```

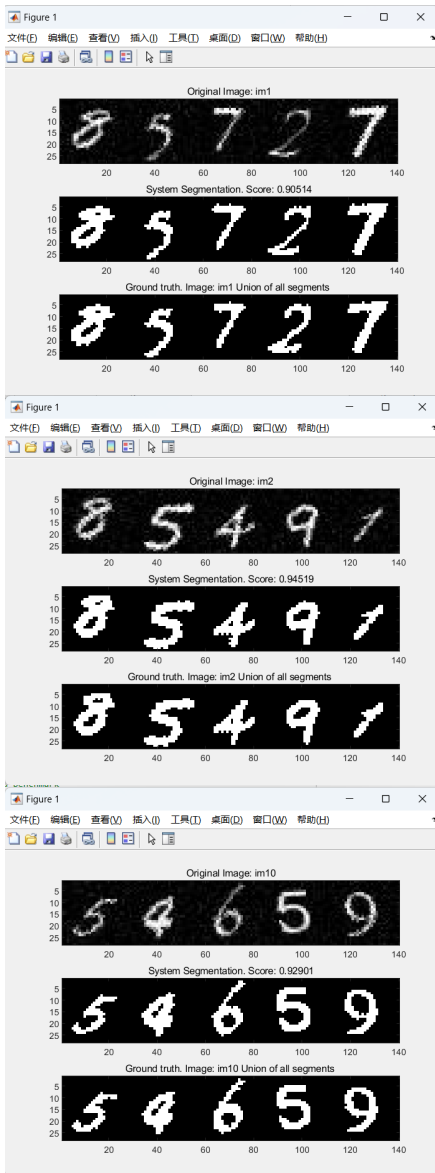
57     if ~merged(i) && ~small(i)
58         for j = i+1:numStats
59             if ~merged(j) && ~small(j)
60                 % Calculate the center coordinates
61                 centro_i = regionprops(S{i}, 'Centroid');
62                 centro_j = regionprops(S{j}, 'Centroid');
63
64                 % Extracting the center coordinates
65                 centro_i = centro_i.Centroid;
66                 centro_j = centro_j.Centroid;
67
68                 % calculate the Euclidean distance
69                 distance = norm(centro_i - centro_j);
70
71                 if distance < Dis_threshold
72                     % Merge two cells and add elements to the
first cell
73                     S{i} = S{i} | S{j};
74                     merged(j) = true;
75                 end
76             end
77         end
78     end
79 end
80
81 S = S(~merged);
82 end

```

You tested 10 images in folder ../datasets/short1
The jaccard scores for all segments in all images were

0.9512	0.8868	0.9302	0.7951	0.9444
0.9010	0.9424	0.8692	0.9658	0.9545
0.9419	0.9454	0.9528	0.9474	0.9333
0.7451	0.9137	0.9438	0.9310	0.9172
0.9262	0.9170	0.9493	0.8958	0.9448
0.9624	0.9298	0.9732	0.9527	0.9931
0.9461	0.8824	0.9115	0.9565	0.9328
0.9298	0.9598	0.7843	0.9282	0.8777
0.9268	0.9203	0.9432	0.9449	0.8901
0.9633	0.7304	0.9231	0.9364	0.7700

The mean of the jaccard scores were 0.91628
This is great!



5 Dimensionality

A:

Dimension k for A:

The set of gray-scale images with 3×2 pixels forms a vector space. To determine the dimension, we need to consider the number of independent basis images that can span this space. In this case, each pixel in a 3×2 image contributes to the dimension, so the total number of pixels is $3 \times 2 = 6$. Therefore, the dimension k for A is 6.

Basis for A:

To define a basis for this vector space, we can choose 6 linearly independent 3×2 images. Here's an example of such a basis:

$$e_1 = \begin{bmatrix} 1 & 0 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}, e_2 = \begin{bmatrix} 0 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}, e_3 = \begin{bmatrix} 0 & 0 \\ 1 & 0 \\ 0 & 0 \end{bmatrix}, e_4 = \begin{bmatrix} 0 & 0 \\ 0 & 1 \\ 0 & 0 \end{bmatrix}, e_5 = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 1 & 0 \end{bmatrix}, e_6 = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 1 \end{bmatrix}$$

Each basis element is a 3×2 image with a single pixel set to 1, and all other pixels set to 0. These basis images are linearly independent and can span the vector space of all 3×2 images.

B:

Dimension k for B:

The set of gray-scale images with 1500×2000 pixels forms a vector space. In this case, the dimension k is equal to the total number of pixels in each image, which is $1500 \times 2000 = 3,000,000$.

Choosing Basis Elements for B:

In the case of such high-dimensional vector spaces, it's impractical to explicitly list individual basis elements. However, you can choose a basis for this space by considering pixel patterns. For example, you can select basis images that represent certain features or patterns commonly found in images. These basis images should be linearly independent and span the entire space.

For instance, you might choose basis images that represent horizontal lines, vertical lines, diagonal lines, gradients, textures, and so on. The choice of basis elements can depend on the specific application or problem you are working on. There are various techniques for automatically extracting or learning basis elements from a set of images, such as Principal Component Analysis (PCA) or Independent Component Analysis (ICA).

The key is to ensure that the chosen basis elements are diverse enough to capture a wide range of image variations and are capable of representing any image in the vector space through linear combinations.

6 Scalar products and norm on images

Scalar Product for Images: The scalar product (or dot product) for images is defined as the sum of the element-wise products of corresponding pixels in two images. If we have two images, u and v , both of the same size, the scalar product $u \cdot v$ is computed as follows:

$$u \cdot v = \sum_{i=1}^M \sum_{j=1}^N u(i, j) \cdot v(i, j)$$

Where M and N are the dimensions (rows and columns) of the images u and v , respectively.

Norm of an Image: The norm of an image represents the "size" or magnitude of the image as if it were a vector. There are various ways to define the norm of an image, but one common approach is to use the Frobenius Norm. For an image u of size $M \times N$, Frobenius Norm, denoted as $\|u\|$, is defined as:

$$\|u\| = \sqrt{\sum_{i=1}^M \sum_{j=1}^N |u(i, j)|^2}$$

```
1  clc
2  clear
3  close all
4  % Given images
5  u = [3 -7; -1 4];
6  v = 1/2 * [1 -1; -1 1];
7  w = 1/2 * [-1 1; -1 1];
8
9  % Calculate norms
10 norm_u = norm(u, 'fro'); % Frobenius norm for the image
11 norm_v = norm(v, 'fro');
12 norm_w = norm(w, 'fro');
13
14 % Calculate scalar products
15 u_dot_v = sum(sum(u .* v));
16 u_dot_w = sum(sum(u .* w));
17 v_dot_w = dot(v, w);
18
19 % Check if matrices u and v_dot_w are orthonormal
20 is_orthonormal = isequal(norm_v, 1) && isequal(norm_w, 1) &&
isequal(dot(v(:), w(:)), 0);
21
22 % Calculate the orthogonal projection of u onto the subspace
spanned by {v, w}
```

```

23 projection = (u_dot_v / (norm_v^2)) * v + (u_dot_w /
24 (norm_w^2)) * w;
25 %%
26 approximation_error = sum(abs(u(:) - projection(:)).^2);
27 u_norm = (norm(u, 'fro'))^2;
28 abs_diff=abs(u(:) - projection(:));
29 diff_norm=(norm(abs_diff(:), 'fro'))^2;
30 diff = diff_norm/u_norm;
31
32 %%
33 % Display results
34 fprintf('Norm of u: %.2f\n', norm_u);
35 fprintf('Norm of v: %.2f\n', norm_v);
36 fprintf('Norm of w: %.2f\n', norm_w);
37 fprintf('Scalar Product u · v: %.2f\n', u_dot_v);
38 fprintf('Scalar Product u · w: %.2f\n', u_dot_w);
39 fprintf('Scalar Product v · w: %.2f\n', v_dot_w);
40 fprintf('Are matrices {v, w} orthonormal? %d\n',
    is_orthonormal);
41 disp('Orthogonal Projection of u onto {v, w}:');
42 disp(projection)
43 disp(['my diff: ', num2str(diff)]);

```

```

Norm of u: 8.66
Norm of v: 1.00
Norm of w: 1.00
Scalar Product u · v: 7.50
Scalar Product u · w: -2.50
Scalar Product v · w: 0.00
Are matrices {v, w} orthonormal? 1
Orthogonal Projection of u onto {v, w}:
    5.0000   -5.0000
   -2.5000    2.5000

my diff: 0.16667
>>

```

Now, let's calculate the norms and scalar products for the given images u, v, and w:

1. $\|u\| : 8.660$
2. $\|v\| : 1$
3. $\|w\| : 1$
4. $u \cdot v = 7.50$
5. $u \cdot w = -2.50$
6. $v \cdot w = 0$

7. Matrices $\{v, w\}$ are orthonormal because their scalar product $(v \cdot w)$ is zero, and the norm of v and w are both one.

8.
$$projection = \begin{bmatrix} 5 & -5 \\ -2.5 & 2.5 \end{bmatrix}$$

9. The projection is the best approximation of u within the subspace

7 Image Compression

1. Background:

- A is a known matrix containing a set of basis vectors as columns.
- x is the parameter vector for which we require a solution, denoting the coefficients of the basis vectors.
- $f(\cdot)$ is the vector form of the observations.

2. **Problem Description:** We wish to find the value of the parameter vector x that best matches the linear model $A * x$ with the observed data $f(\cdot)$.

3. Objective of least squares:

Minimize the norm of the residual vector, i.e., minimize the following equation:

$$minimize ||A * x - f(\cdot)||_2^2$$

This is equivalent to finding the parameter vector x such that the linear model $A * x$ is as close as possible to the observed data $f(\cdot)$.

4. Solution process:

- By computing $A * x$, we can obtain an estimate of the linear model.
- Calculate the residual vector: $residual = A * x - f(\cdot)$.
- The goal of least squares is to find the parameter vector x that minimizes the norm of the residual vector $residual$.
- This is accomplished by solving the following regular equation:
$$A' * A * x = A' * f(\cdot)$$

where A' denotes the transpose matrix of A .
- Ultimately, the x -value will be solved such that $A * x$ is closest to $f(\cdot)$ and the residual vector $residual$ minimizes the norm.

5. Solve for the value of x :

- It is convenient to solve regular equations to find the value of x using the left division operator $()$ in MATLAB. Specifically, $x = A \setminus f(\cdot)$ will automatically compute the regular equation and solve for the value of x .

```
1 | c1c
2 | clear
```

```

3 close all
4
5 % Define the basis images
6 phi1 = 1/2 * [1 0 -1; 1 0 -1; 0 0 0; 0 0 0];
7 phi2 = 1/3 * [1 1 1; 1 0 1; -1 -1 -1; 0 -1 0];
8 phi3 = 1/3 * [0 1 0; 1 1 1; 1 0 1; 1 1 1];
9 phi4 = 1/2 * [0 0 0; 0 0 0; 1 0 -1; 1 0 -1];
10
11
12 % Define the original image f
13 f = [-2 6 3; 13 7 5; 7 1 8; -3 4 4];
14
15 % Verify orthonormality of basis images
16 orthonormality = isequal(norm(phi1,1),1) &&
isequal(norm(phi2,1),1) && ...
17 isequal(norm(phi3,1),1) && isequal(norm(phi4,1),1) &&...
18 isequal(dot(phi3(:), phi4(:)), 0) && ...
19 isequal(dot(phi1(:), phi2(:)), zeros(size(3))) && ...
20 isequal(dot(phi1(:), phi3(:)), zeros(size(3))) && ...
21 isequal(dot(phi1(:), phi4(:)), zeros(size(3))) && ...
22 isequal(dot(phi2(:), phi3(:)), zeros(size(3))) && ...
23 isequal(dot(phi2(:), phi4(:)), zeros(size(3)));
24
25
26 %% pseudo-inverse
27 % % Stack the basis images into a matrix
28 % A = [phi1(:), phi2(:), phi3(:), phi4(:)];
29 %
30 % % Calculate the coefficients using the pseudo-inverse
31 % x = pinv(A) * f(:);
32 %
33 % % Reconstruct the approximate image
34 % fa = A * x;
35 %
36 %
37 % fa_matrix = reshape(fa, 4, 3);
38
39 %%
40 % Stack the basis images into a matrix
41 A = [phi1(:), phi2(:), phi3(:), phi4(:)];
42
43 % Calculate the coefficients using the
44 x = A \ f(:);
45
46 % Reconstruct the approximate image
47 fa = A * x;
48
49 % Calculate the approximation error
50 approximation_error = sum(abs(f(:) - fa).^2);
51 % approximation_error = norm(f(:) - fa,'fro');
52

```

```

53 fa_matrix = reshape(fa, 4, 3);
54 %%
55 % Display results
56 disp('Orthonormality of Basis Images:');
57 disp(['Are basis images orthonormal ? orthonormality = ',
num2str(orthonormality)]);
58 disp(['Coordinates (x1, x2, x3, x4):']);
59 disp(x);
60 disp('Approximate Image fa:');
61 disp(fa_matrix);
62 disp(['Norm Approximation Error: ',
num2str(approximation_error)]);

```

```

Orthonormality of Basis Images:
Are basis images orthonormal ? orthonormality = 1
Coordinates (x1, x2, x3, x4):
    1.5000
    1.6667
   17.0000
   -4.0000

Approximate Image fa:
    1.3056    6.2222   -0.1944
    6.9722    5.6667    5.4722
    3.1111   -0.5556    7.1111
    3.6667    5.1111    7.6667

Norm Approximation Error: 136.9722
my diff: 0.30643

```

I think the result of task 2 is better, considering that there is a difference in the dimension of the matrix, I first calculate the norm of the error between the elements of the matrix, and then calculate the ratio of diff_norm and Norm Approximation Error as a basis for judgment, the smaller the ratio the more approximate it is.

8 Image Bases

```

1  clc
2  clear
3  close all
4  % Load the dataset
5  load('assignment1bases.mat');
6
7  % Initialize variables to store the mean error norms
8  mean_error_norms = zeros(2, 3); % Rows: test sets, Columns:
bases
9
10 %% Iterate through the test sets (general and face)

```

```

11 for test_set = 1:2
12     % Select the test images from the corresponding stack
13     test_images = stacks{test_set};
14
15     % Iterate through the three bases
16     for basis_idx = 1:3
17         % Select the basis for this iteration
18         basis = bases{basis_idx};
19
20         % Initialize an array to store error norms for
individual images
21         error_norms = zeros(size(test_images, 3), 1);
22         % Iterate through all test images
23         for img_idx = 1:400
24             % Get the current test image
25             img = test_images(:, :, img_idx);
26
27             % Project the image onto the basis and calculate
the error norm
28             [up, r] = projectAndCalculateError(img, basis);
29
30             % Store the error norm
31             error_norms(img_idx) = r;
32         end
33
34         % Calculate the mean error norm for this basis and
test set
35         mean_error_norm = mean(error_norms);
36
37         % Store the result in the mean_error_norms matrix
38         mean_error_norms(test_set, basis_idx) =
mean_error_norm;
39
40     end
41 end
42
43 %% print result
44 % choose the test image
45 plotset_idx=2;
46 plotting_idx=375;
47 test_images = stacks{plotset_idx};
48 plot_img = zeros(19,19,3);
49 for basis_idx = 1:3
50     basis = bases{basis_idx};
51     [plot_up, r] = projectAndCalculateError(test_images(:,
:, plotting_idx), basis);
52     plot_img(:,:,basis_idx)=plot_up;
53 end
54
55 % plot results images
56 figure;

```



```

57 subplot(1, 4, 1);
58 imshow(uint8(test_images(:, :, plotimg_idx)));
59 title('Test Image');
60
61 subplot(1, 4, 2);
62 imshow(uint8(plot_img(:, :, 1)));
63 title('Projection1');
64
65 subplot(1, 4, 3);
66 imshow(uint8(plot_img(:, :, 2)));
67 title('Projection2');
68
69 subplot(1, 4, 4);
70 imshow(uint8(plot_img(:, :, 3)));
71 title('Projection3');
72
73 basis1 = abs(bases{1});
74 min_value = min(basis1(:));
75 max_value = max(basis1(:));
76 basis1 = 255 * (basis1 - min_value) / (max_value - min_value);
77 basis1 = uint8(basis1);
78
79 figure;
80 subplot(1, 4, 1);
81 imshow(basis1(:, :, 1));
82 title('Basis1');
83 subplot(1, 4, 2);
84 imshow(basis1(:, :, 2));
85 subplot(1, 4, 3);
86 imshow(basis1(:, :, 3));
87 subplot(1, 4, 4);
88 imshow(basis1(:, :, 4));
89
90 % Display or use the mean_error_norms matrix as needed
91 disp('Mean Error Norms:');
92 disp(mean_error_norms);


```

```


1 function [up, r] = projectAndCalculateError(u, basis)
2     % Flatten the image into a column vector
3     reshape_u = u(:) ;
4     % Create a matrix containing the basis vectors as columns
5     reshape_basis = reshape(basis, [], 4);
6     x = reshape_basis \ u(:);
7     up = reshape_basis * x;
8     % Calculate the error norm
9     r = norm(u(:) - up, "fro");
10    % r = sum(abs(reshape_u - up).^2);
11    up = reshape(up, 19, 19);
12 end
13

```

Basis1:







Basis2:







Basis3:



- TestSet1 :



- TestSet2 :



- Mean Error Norms:

Mean Error Norms:

649.2013	795.1902	697.3214
860.4754	821.0271	944.9009

Basis 1 works best on test set 1 ,because the mean of the error norm on test set 1 is the smallest.

And basis 2 works best on test set 2 ,because the mean of the error norm on test set 1 is the smallest. Moreover, test set 2 is all face images, and the images processed by basis2 are also face images.